

In [1]:

```
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_log_error
from statsmodels.tsa.deterministic import DeterministicProcess
```

In [2]:

```
# path to the dataset in Kaggle's notebook
path = '../input/store-sales-time-series-forecasting/'
```

1. Compute Moving Average of Oil Prices

In [3]:

```
# read oil price
data_oil = pd.read_csv(path + 'oil.csv', parse_dates=['date'], infer_datetime_format=True, index_col='date')

#####
# TODO: compute data_oil['ma_oil'] as the moving average of data_oil['dcoilwtico'] with window size 7
# Hint: check the documentation of .rolling() method of pandas.DataFrame
#####
data_oil['ma_oil'] = data_oil['dcoilwtico'].rolling(7).mean()

# Create contiguous moving average of oil prices
calendar = pd.DataFrame(index=pd.date_range('2013-01-01', '2017-08-31'))

#####
# TODO 1: merge two DataFrame instances (data_oil and calendar) such that the merged instances has the same indexes
# as calendar.
# TODO 2: replace each NaN in data_oil['ma_oil'] by the first non-null value before it.
# Hint: check the documentation of .merge() and .fillna() methods of pandas.DataFrame
#####
calendar = pd.merge(calendar, data_oil, left_index=True, how='left', right_index=True)
calendar.fillna(method='bfill', inplace=True)
calendar.head(15) # display some entries of calendar
```

Out[3]:

	dcoilwtico	ma_oil
2013-01-01	93.14	93.218571
2013-01-02	93.14	93.218571
2013-01-03	92.97	93.218571
2013-01-04	93.12	93.218571
2013-01-05	93.20	93.218571
2013-01-06	93.20	93.218571
2013-01-07	93.20	93.218571
2013-01-08	93.21	93.218571
2013-01-09	93.08	93.218571
2013-01-10	93.81	93.218571
2013-01-11	93.60	93.284286
2013-01-12	94.27	93.470000
2013-01-13	94.27	93.470000
2013-01-14	94.27	93.470000
2013-01-15	93.26	93.490000

2. Create Workday Feature

In [4]:

```
#####
# TODO: create a True/False feature calendar['wd'] to indicate whether each date is a workday (Monday-Friday) or not.
# Hint: check documentation of pandas.DatetimeIndex.dayofweek
#####
calendar['wd'] = calendar.index.day_of_week <=4

calendar.head(15) # display some entries of calendar
```

Out[4]:

	dcoilwtico	ma_oil	wd
2013-01-01	93.14	93.218571	True
2013-01-02	93.14	93.218571	True
2013-01-03	92.97	93.218571	True
2013-01-04	93.12	93.218571	True
2013-01-05	93.20	93.218571	False
2013-01-06	93.20	93.218571	False
2013-01-07	93.20	93.218571	True
2013-01-08	93.21	93.218571	True
2013-01-09	93.08	93.218571	True
2013-01-10	93.81	93.218571	True
2013-01-11	93.60	93.284286	True
2013-01-12	94.27	93.470000	False
2013-01-13	94.27	93.470000	False
2013-01-14	94.27	93.470000	True
2013-01-15	93.26	93.490000	True

3. Read Train and Test Data

In [5]:

```
df_train = pd.read_csv(path + 'train.csv',
                        usecols=['store_nbr', 'family', 'date', 'sales'],
                        dtype={'store_nbr': 'category', 'family': 'category', 'sales': 'float32'},
                        parse_dates=['date'], infer_datetime_format=True)

df_train.date = df_train.date.dt.to_period('D')
df_train = df_train.set_index(['store_nbr', 'family', 'date']).sort_index()

df_train.head(15) # display some entries of the training data
```

Out[5]:

		sales	
store_nbr	family	date	
1	AUTOMOTIVE	2013-01-01	0.0
		2013-01-02	2.0
		2013-01-03	3.0
		2013-01-04	3.0
		2013-01-05	5.0
		2013-01-06	2.0
		2013-01-07	0.0
		2013-01-08	2.0
		2013-01-09	2.0
		2013-01-10	2.0
		2013-01-11	3.0
		2013-01-12	2.0
		2013-01-13	2.0
		2013-01-14	2.0
		2013-01-15	1.0

In [6]:

```
df_test = pd.read_csv(path + 'test.csv',
                      usecols=['store_nbr', 'family', 'date'],
                      dtype={'store_nbr': 'category', 'family': 'category'},
                      parse_dates=['date'], infer_datetime_format=True)

df_test.date = df_test.date.dt.to_period('D')
df_test = df_test.set_index(['store_nbr', 'family', 'date']).sort_index()

df_test.head(15) # display some entries of the testing data
```

Out[6]:

	store_nbr	family	date
			2017-08-16
			2017-08-17
			2017-08-18
			2017-08-19
			2017-08-20
			2017-08-21
			2017-08-22
1	AUTOMOTIVE		2017-08-23
			2017-08-24
			2017-08-25
			2017-08-26
			2017-08-27
			2017-08-28
			2017-08-29
			2017-08-30

In [7]:

```
# set the range of data used in training
sdate = '2017-04-01'
edate = '2017-08-15'

# we will train a model that takes feature of a date as input and predicts the sales for each store and family of goods on that d
y = df_train.unstack(['store_nbr', 'family']).loc[sdate:edate]

#####
# TODO: create the trend feature X: the value for sdate is 1, the value for the next day of sdate is 2, etc.
# Hint: check the documentation of DeterministicProcess, or this tutorial: https://www.kaggle.com/code/ryanholbrook/trend.
#####
dp = DeterministicProcess(
    index=y.index, # dates from the training data
    constant=True, # dummy feature for the bias (y_intercept)
    order=1,       # the time dummy (trend)
    drop=True,     # drop terms if necessary to avoid collinearity
)
X = dp.in_sample()

# Extentions
X['oil'] = calendar.loc[sdate:edate]['ma_oil'].values
X['wd'] = calendar.loc[sdate:edate]['wd'].values

X.head(15)
```

Out[7]:

	const	trend	oil	wd
date				
2017-04-01	1.0	1.0	49.034286	False
2017-04-02	1.0	2.0	49.034286	False
2017-04-03	1.0	3.0	49.034286	True
2017-04-04	1.0	4.0	49.561429	True
2017-04-05	1.0	5.0	50.150000	True
2017-04-06	1.0	6.0	50.625714	True
2017-04-07	1.0	7.0	51.022857	True
2017-04-08	1.0	8.0	51.417143	False
2017-04-09	1.0	9.0	51.417143	False
2017-04-10	1.0	10.0	51.417143	True
2017-04-11	1.0	11.0	51.822857	True
2017-04-12	1.0	12.0	52.232857	True
2017-04-13	1.0	13.0	52.547143	True
2017-04-14	1.0	14.0	50.512857	True
2017-04-15	1.0	15.0	50.512857	False

4. Train Model!

In [8]:

```
model = LinearRegression()
model.fit(X, y)
y_pred = pd.DataFrame(model.predict(X), index=X.index, columns=y.columns)
```

In [9]:

```
# Results on the training set

y_pred = y_pred.stack(['store_nbr', 'family']).reset_index()
y_target = y.stack(['store_nbr', 'family']).reset_index().copy()

y_target['sales_pred'] = y_pred['sales'].clip(0.) # Sales should be >= 0

#####
# TODO: show the training loss for each type of product.
# Hint: check the documentation of DataFrame.groupby() and GroupBy.apply().
#####
```

In [10]:

```
y_target.groupby('family')[['sales', 'sales_pred']].apply(lambda Y: mean_squared_log_error(Y.sales, Y.sales_pred))
```

Out[10]:

family	
AUTOMOTIVE	0.274819
BABY CARE	0.070619
BEAUTY	0.284931
BEVERAGES	0.203731
BOOKS	0.028635
BREAD/BAKERY	0.135292
CELEBRATION	0.339155
CLEANING	0.212288
DAIRY	0.147564
DELI	0.119779
EGGS	0.202576
FROZEN FOODS	0.185090
GROCERY I	0.214652
GROCERY II	0.371296
HARDWARE	0.291139
HOME AND KITCHEN I	0.278487
HOME AND KITCHEN II	0.232872
HOME APPLIANCES	0.162094
HOME CARE	0.131746
LADIESWEAR	0.277219
LAWN AND GARDEN	0.295717
LINGERIE	0.421013
LIQUOR,WINE,BEER	0.791481
MAGAZINES	0.274324
MEATS	0.174464
PERSONAL CARE	0.150410
PET SUPPLIES	0.222304
PLAYERS AND ELECTRONICS	0.237393
POULTRY	0.156695
PREPARED FOODS	0.136247
PRODUCE	0.229127
SCHOOL AND OFFICE SUPPLIES	1.388927
SEAFOOD	0.286115

dtype: float64

In [11]:

```
# Test predictions

stest = '2017-08-16'
etest = '2017-08-31'

#####
# TODO: create the feature matrix of test data.
# Hint: check the documentation of DeterministicProcess.
#####
# X_test = None # change 'None' to your answer
# dp = DeterministicProcess(
#     index=df_test.unstack(['store_nbr', 'family']).index, # dates from the training data
#     constant=True,      # dummy feature for the bias (y_intercept)
#     order=1,            # the time dummy (trend)
#     drop=True,          # drop terms if necessary to avoid collinearity
# )
# X_test = dp.in_sample()
X_test = dp.out_of_sample(steps=16)
# Extensions
X_test['oil'] = calendar.loc[stest:etest]['ma_oil'].values
X_test['wd'] = calendar.loc[stest:etest]['wd'].values

print(X_test)

sales_pred = pd.DataFrame(model.predict(X_test), index=X_test.index, columns=y.columns)
sales_pred = sales_pred.stack(['store_nbr', 'family'])

sales_pred[sales_pred < 0] = 0. # Sales should be >= 0
```

	const	trend	oil	wd
2017-08-16	1.0	138.0	48.281429	True
2017-08-17	1.0	139.0	47.995714	True
2017-08-18	1.0	140.0	47.852857	True
2017-08-19	1.0	141.0	47.688571	False
2017-08-20	1.0	142.0	47.688571	False
2017-08-21	1.0	143.0	47.688571	True
2017-08-22	1.0	144.0	47.522857	True
2017-08-23	1.0	145.0	47.645714	True
2017-08-24	1.0	146.0	47.598571	True
2017-08-25	1.0	147.0	47.720000	True
2017-08-26	1.0	148.0	47.624286	False
2017-08-27	1.0	149.0	47.624286	False
2017-08-28	1.0	150.0	47.624286	True
2017-08-29	1.0	151.0	47.320000	True
2017-08-30	1.0	152.0	47.115714	True
2017-08-31	1.0	153.0	47.060000	True

In [12]:

```
sales_pred
```

Out[12]:

		sales	
store_nbr		family	
2017-08-16	1	AUTOMOTIVE	4.922374
		BABY CARE	0.000000
		BEAUTY	3.543350
		BEVERAGES	2192.004787
		BOOKS	0.140162
...
2017-08-31	9	POULTRY	377.218940
		PREPARED FOODS	91.890624
		PRODUCE	1477.958478
		SCHOOL AND OFFICE SUPPLIES	100.100853
		SEAFOOD	14.306982

28512 rows × 1 columns

In [13]:

```
# Create submission
```

```
df_sub = pd.read_csv(path + 'sample_submission.csv', index_col='id')  
df_sub.sales = sales_pred.values  
df_sub.to_csv('submission.csv', index=True)
```

In []: